



SEARCH TREES

School of Artificial Intelligence


LAST LECTURE:

(2,4) TREES AND RED-BLACK TREES

- Multi-way search tree: Internal nodes may have more than two children
 - let w be a node of an ordered tree
 - w is a d -node if w has d children
 - Each internal node of T has at least two children. That is, each internal node is a d -node such that $d \geq 2$.
 - Each internal d -node w of T with children c_1, \dots, c_d stores an ordered set of $d - 1$ key-value pairs $(k_1, v_1), \dots, (k_{d-1}, v_{d-1})$, where $k_1 \leq \dots \leq k_{d-1}$.
 - Let us conventionally define $k_0 = -\infty$ and $k_d = +\infty$. For each item (k, v) stored at a node in the subtree of w rooted at c_i , $i = 1, \dots, d$, we have that $k_{i-1} \leq k \leq k_i$.
- A d -node stores $d-1$ regular keys
- External nodes do not store any data and serve only as “placeholders”
 - Reference to None
- An n -item multiway search tree has $n+1$ external nodes

MULTIWAY SEARCH TREE

- Search:
 - Start from the root
 - Compare k with d-nodes
 - Search successful
 - Locate key
 - Search unsuccessful
 - Reach an external node




DATA STRUCTURE FOR MULTIWAY SEARCH TREES

- General tree: linked data structure
- Secondary container: finding the smallest key at the node that is greater than or equal to k
 - Sorted map: `find_ge(k)`
 - `SortedTableMap` from previous lecture
 - Associated value in case of a match for key k , or the child c_i such that $k_{i-1} < k < k_i$
 - k_i in the secondary structure to pair (v_i, c_i)
 - Process d -node when searching for an item of T with key k can be performed using binary search in $O(\log d)$, d = number of children
 - d_{\max} = maximum number of children of any node of T , h = height of T , search time in a multiway search tree is $O(h \log d_{\max})$
 - If d_{\max} is a constant, the running time for performing a search is $O(h)$

(2,4) TREES

- Sometimes 2-4 tree or 2-3-4 tree
- **Size property:** every internal node has at most four children
- **Depth property:** all external nodes have the same depth
- The height of a 2-4 tree storing n items is $O(\log n)$
 - Sufficient to keep the tree balanced
 - Search takes $O(\log n)$ time




(2,4) TREES INSERTION


- Insert new item (k, v) , search for k
 - k not in T : locate an external node z .
 - if w is the parent of z .
 - insert new item into node w and add a new child y to w on the left of z
 - may violate the size property
 - 4-node becomes 5-node after the insertion – **overflow**
 - Resolution: **split**
 - Replace w with two nodes w' and w'' , where
 - w' is a 3-node with children c_1, c_2, c_3 storing keys k_1 and k_2
 - w'' is a 2-node with children c_4, c_5 storing key k_4
 - If w is the root of T , create a new root node u ; else, let u be the parent of w
 - Insert key k_3 into u and make w' and w'' children of u , so that if w was child i of u , then w' and w'' become children i and $i+1$ of u

(2,4) TREES INSERTION


- Node split



(a)







(b)



(c)



(2,4) TREES INSERTION

- Node split



(2,4) TREES INSERTION

- Node split



(2,4) TREES INSERTION

- Analysis
 - d_{\max} is at most 4, search for the placement of new key k uses $O(1)$ time at each level, and $O(\log n)$ time overall
 - Split operations: bounded by the height of the tree
 - Insertion process runs in $O(\log n)$ time

(2,4) TREES DELETION


- Removal of an item
 - Search for k in T
 - Can always be reduced to - the item to be removed is stored at a node w whose children are external nodes
 - Item with key k to be removed is stored in the i^{th} item (k_i, v_i) at a node z that has only internal-node children,
 - swap item (k_i, v_i) with an appropriate item that is stored at a node w with external-node children
 - Find the rightmost internal node w in the subtree rooted at the i^{th} child of z
 - Swap the item (k_i, v_i) at z with the last item of w

(2,4) TREES DELETION


- Removal preserves the depth property, may violate the size property at w
 - 2-node becomes a 1-node with no items at all – **underflow**
 - Check if an immediate sibling s of w is a 3-node or a 4-node, then perform a **transfer**: move a child of s to w , a key of s to the parent u of w and s , and a key of u to w
 - If w has only one sibling, or if both immediate siblings of w are 2-nodes, then perform a **fusion**: merge w with a sibling to a new node w' and move a key from the parent u of w to w'

(2,4) TREES DELETION


- Fusion at node w may cause a new underflow to occur at the parent u of w , which triggers a transfer or fusion at u
- Number of fusion operations is bounded by the height of the tree – $O(\log n)$




(a)



(b)



(c)



(d)

(2,4) TREES


PERFORMANCE

- Identical to AVL tree
 - Height of a 2-4 tree with n items is $O(\log n)$
 - Split, transfer, fusion: $O(1)$
 - Search, insertion, removal: $O(\log n)$

Operation	Running Time
$k \text{ in } T$	$O(\log n)$
$T[k] = v$	$O(\log n)$
$T.\text{delete}(p), \text{del } T[k]$	$O(\log n)$
$T.\text{find_position}(k)$	$O(\log n)$
$T.\text{first}(), T.\text{last}(), T.\text{find_min}(), T.\text{find_max}()$	$O(\log n)$
$T.\text{before}(p), T.\text{after}(p)$	$O(\log n)$
$T.\text{find_lt}(k), T.\text{find_le}(k), T.\text{find_gt}(k), T.\text{find_ge}(k)$	$O(\log n)$
$T.\text{find_range}(\text{start}, \text{stop})$	$O(s + \log n)$
$\text{iter}(T), \text{reversed}(T)$	$O(n)$


RED-BLACK TREES

- AVL trees – need to perform rotations
- 2-4 trees – need to perform split and fusion operations
- Red-black trees: $O(1)$ structural changes after an update to stay balanced
- Red-black tree
 - Binary search tree, nodes coloured
 - **Root property:** root is black
 - **Red property:** the children of a red node are black
 - **Depth property:** all nodes with zero or one children have the same black depth
 - **Black depth:** number of black ancestors



RED-BLACK TREES

- Red-black trees and 2-4 trees
 - Given a red-black tree, a 2-4 tree can be constructed by merging every red node w into its parent,
 - storing the entry from w at its parent,
 - and with the children of w becoming ordered children of the parent
- Height of a red-black tree – $O(\log n)$






RED-BLACK TREES INSERTION

- Search: similar to binary search tree - $O(\log n)$
 - returns position x
- If x is the root, colour it black
- Other cases, colour x red
- Insertion preserves the root and depth properties, but may violate the red property
- if x is not the root of T and parent y of x is red – double red situation


RED-BLACK TREES

INSERTION

- Case 1: The sibling s of y is black (or None)
- Trinode restructuring
 - Node x, y, z
 - Label them a, b , and c
 - Replace z with the node labeled b and make nodes a and c the children of b
 - Colour b black and a and c red



(a)






(b)


RED-BLACK TREES

INSERTION

- Case 2: sibling s of y is red
- Overflow in its equivalent 2-4 tree
- Fix: **split/recolouring**
- Colour y and s black and their parent z red
- If z is root, it remains black
 - Unless z is the root, the portion of any path through the affected part of the tree is incident to one black node



(a)




(b)

RED-BLACK TREES


INSERTION




(a)




(b)




(c)




(d)




(e)



(f)




(g)




(h)

RED-BLACK TREES


INSERTION




(i)



(j)




(k)




(l)


RED-BLACK TREES INSERTION




(m)



(n)




(o)



(p)

RED-BLACK TREES


INSERTION



RED-BLACK TREES





DELETION

- Search – $O(\log n)$
- If removed node is red – no affect on the black depth property, or red violations
- If removed node is black and has one child that was a red leaf
 - Recolour solves the problem
- If removed node is a black leaf
 - Black deficit of 1
 - Removed node must have a sibling whose subtree has black height 1
 - More general setting with a node z with two subtrees: T_{heavy} and T_{light} . Black depth of T_{heavy} is one more than T_{light}
 - Z : parent of removed leaf
 - Y : root of T_{heavy}





RED-BLACK TREES DELETION

- Case 1: node y is black and has a red child x
- Trinode restructuring:
- x, y and z
- a, b and c
- Make b the parent of the other two
- Colour a and c black
- Give b the previous colour of z





RED-BLACK TREES DELETION

- Case 2: node y is black and both children of y are black (or None)
- Recolouring: colour y red and if z is red, colour it black
- Z becomes deficient, repeat consideration of all three cases at the parent of z as a remedy



(a)




(b)

RED-BLACK TREES


DELETION

- Case 3: node y is red
- Rotation about y and z
- Recolor y black and z red
- Repeat step 1, 2 and 3 if necessary




RED-BLACK TREES


DELETION




(a)



(b)




(c)




(d)

RED-BLACK TREES


DELETION




(e)




(f)




(g)




(h)



(i)



(j)



(k)

RED BLACK TREES

PERFORMANCE

- Identical to AVL tree
 - Height of a red-black tree is $O(\log n)$
 - Insertion requires $O(\log n)$ recolourings and at most one trinode restructuring
 - Deletion requires $O(\log n)$ reclourings and at most two restructuring operations

Operation	Running Time
$k \text{ in } T$	$O(\log n)$
$T[k] = v$	$O(\log n)$
$T.\text{delete}(p), \text{del } T[k]$	$O(\log n)$
$T.\text{find_position}(k)$	$O(\log n)$
$T.\text{first}(), T.\text{last}(), T.\text{find_min}(), T.\text{find_max}()$	$O(\log n)$
$T.\text{before}(p), T.\text{after}(p)$	$O(\log n)$
$T.\text{find_lt}(k), T.\text{find_le}(k), T.\text{find_gt}(k), T.\text{find_ge}(k)$	$O(\log n)$
$T.\text{find_range}(\text{start}, \text{stop})$	$O(s + \log n)$
$\text{iter}(T), \text{reversed}(T)$	$O(n)$

PYTHON IMPLEMENTATION

```
1 class RedBlackTreeMap(TreeMap):
2     """Sorted map implementation using a red-black tree."""
3     class _Node(TreeMap._Node):
4         """Node class for red-black tree maintains bit that denotes color."""
5         __slots__ = '_red'    # add additional data member to the Node class
6
7     def __init__(self, element, parent=None, left=None, right=None):
8         super().__init__(element, parent, left, right)
9         self._red = True    # new node red by default
10
11     #----- positional-based utility methods -----
12     # we consider a nonexistent child to be trivially black
13     def _set_red(self, p): p._node._red = True
14     def _set_black(self, p): p._node._red = False
15     def _set_color(self, p, make_red): p._node._red = make_red
16     def _is_red(self, p): return p is not None and p._node._red
17     def _is_red_leaf(self, p): return self._is_red(p) and self.is_leaf(p)
18
19     def _get_red_child(self, p):
20         """Return a red child of p (or None if no such child)."""
21         for child in (self.left(p), self.right(p)):
22             if self._is_red(child):
23                 return child
24         return None
```

PYTHON IMPLEMENTATION

```
26 def _rebalance_insert(self, p):
27     self._resolve_red(p)                # new node is always red
28
29 def _resolve_red(self, p):
30     if self.is_root(p):
31         self._set_black(p)              # make root black
32     else:
33         parent = self.parent(p)
34         if self._is_red(parent):          # double red problem
35             uncle = self.sibling(parent)
36             if not self._is_red(uncle):    # Case 1: misshapen 4-node
37                 middle = self._restructure(p) # do trinode restructuring
38                 self._set_black(middle)    # and then fix colors
39                 self._set_red(self.left(middle))
40                 self._set_red(self.right(middle))
41             else:                          # Case 2: overfull 5-node
42                 grand = self.parent(parent)
43                 self._set_red(grand)        # grandparent becomes red
44                 self._set_black(self.left(grand)) # its children become black
45                 self._set_black(self.right(grand))
46                 self._resolve_red(grand)    # recur at red grandparent
```

```
48 def _rebalance_delete(self, p):
49     if len(self) == 1:
50         self._set_black(self.root())    # special case: ensure that root
51     elif p is not None:
52         n = self.num_children(p)
53         if n == 1:                       # deficit exists unless child is a
54             c = next(self.children(p))
55             if not self._is_red_leaf(c):
56                 self._fix_deficit(p, c)
57         elif n == 2:                     # removed black node with red
58             if self._is_red_leaf(self.left(p)):
59                 self._set_black(self.left(p))
60             else:
61                 self._set_black(self.right(p))
```


PYTHON IMPLEMENTATION

```
63 def _fix_deficit(self, z, y):
64     """ Resolve black deficit at z, where y is the root of z's heavier subtree. """
65     if not self._is_red(y): # y is black; will apply Case 1 or 2
66         x = self._get_red_child(y)
67         if x is not None: # Case 1: y is black and has red child x; do "transfer"
68             old_color = self._is_red(z)
69             middle = self._restructure(x)
70             self._set_color(middle, old_color) # middle gets old color of z
71             self._set_black(self.left(middle)) # children become black
72             self._set_black(self.right(middle))
73         else: # Case 2: y is black, but no red children; recolor as "fusion"
74             self._set_red(y)
75             if self._is_red(z):
76                 self._set_black(z) # this resolves the problem
77             elif not self.is_root(z):
78                 self._fix_deficit(self.parent(z), self.sibling(z)) # recur upward
79         else: # Case 3: y is red; rotate misaligned 3-node and repeat
80             self._rotate(y)
81             self._set_black(y)
82             self._set_red(z)
83             if z == self.right(y):
84                 self._fix_deficit(z, self.left(z))
85             else:
86                 self._fix_deficit(z, self.right(z))
```

PREVIOUSLY ON BINARY SEARCH TREES

- Binary Search Tree
 - Performance: $O(h)$
- Balancing search tree
 - Rotation
 - X-Y rotation
 - Trinode rotation
- AVL tree
 - Height of AVL tree: number of nodes in a path
 - Height balance property
- Splay tree
 - Splay operations: search, add, remove

NAVIGATING A BINARY SEARCH TREE

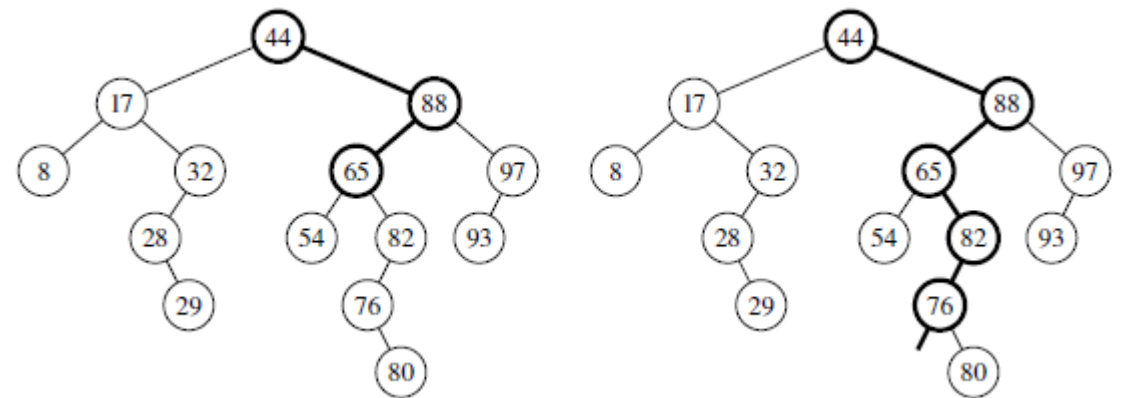
- In order traversal of a binary search tree visits positions in increasing order of their keys
- Proof by induction
 - Base: tree has one item
 - Inductive: recursive inorder traversal: left child(ren) \rightarrow node \rightarrow right child(ren), by binary search tree property, inorder traversal visits positions in increasing order
- Inorder traversal: $O(n) \Rightarrow$ sorted iteration of the keys of a map in $O(n)$, provided that the map is represented as a binary search tree

BINARY SEARCH TREE ADT

- `first()`: returns the position containing the least key, or `None` if the tree is empty
- `last()`: returns the position containing the greatest key, or `None` if empty tree
- `before(p)`: returns the position containing the greatest key that is less than that of position `p`, or `None` if `p` is the first position
- `after(p)`: returns position containing the least key that is greater than that of the position `p`, or `None` if `p` is the last position

SEARCHES

- Locate a particular key by viewing it as a decision tree
- At each position p : is the desired k less than, equal to, or greater than the key stored at position p ?



```
Algorithm TreeSearch( $T, p, k$ ):  
    if  $k == p.key()$  then  
        return  $p$   
    else if  $k < p.key()$  and  $T.left(p)$  is not None then  
        return TreeSearch( $T, T.left(p), k$ )  
    else if  $k > p.key()$  and  $T.right(p)$  is not None then  
        return TreeSearch( $T, T.right(p), k$ )  
    return  $p$ 
```

INSERTIONS

- Map backed by a binary search tree
- $M[k] = v$
 - Search for key k
 - If found, update value
 - Otherwise, create a new node and insert it into the binary search tree
- E.g. insert 68 into tree

Algorithm TreeInsert(T, k, v):

Input: A search key k to be associated with value v

$p = \text{TreeSearch}(T, T.\text{root}(), k)$

if $k == p.\text{key}()$ **then**


Set p 's value to v

else if $k < p.\text{key}()$ **then**

add node with item (k, v) as left child of p

else

add node with item (k, v) as right child of p


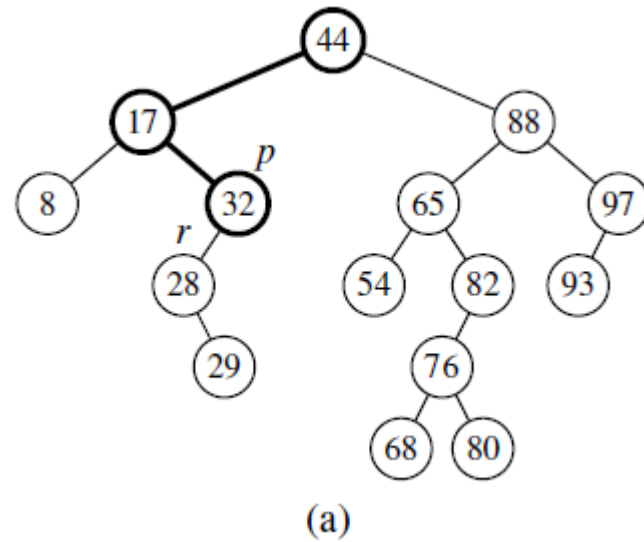


DELETION

- Find the position p of T storing an item with key equal to k , if the search is successful:
 1. If p has at most one child, then delete p , replace it with the child
 2. If p has two children
 - Locate position r , where $r = \text{before}(p)$. r is the rightmost position of the left subtree of p
 - Use r 's item as a replacement for position p
 - Delete node at r , since r has at most 1 child, repeat step 1 for r


DELETION

- Delete item with $k=32$ with one child r



DELETION

- Delete item with $k=88$




PERFORMANCE OF A BINARY SEARCH TREE

- Almost all operations have a worst-case running time of $O(h)$
- Single call to `after()` is worst case $O(h)$, n successive calls made during a call to `__iter__` require a total of $O(n)$ time" each edge is traced at most twice
- $O(1)$ amortised time bounds
- Is $O(h)$ same as $O(\log n)$?
- No, BST can be unbalanced

Operation	Running Time
$k \text{ in } T$	$O(h)$
$T[k], T[k] = v$	$O(h)$
$T.\text{delete}(p), \text{del } T[k]$	$O(h)$
$T.\text{find_position}(k)$	$O(h)$
$T.\text{first}(), T.\text{last}(), T.\text{find_min}(), T.\text{find_max}()$	$O(h)$
$T.\text{before}(p), T.\text{after}(p)$	$O(h)$
$T.\text{find_lt}(k), T.\text{find_le}(k), T.\text{find_gt}(k), T.\text{find_ge}(k)$	$O(h)$
$T.\text{find_range}(\text{start}, \text{stop})$	$O(s + h)$
$\text{iter}(T), \text{reversed}(T)$	$O(n)$


BALANCED SEARCH TREES

- Balanced binary search tree: $O(\log n)$ time for basic map operations
- What about the running time of operations after some sequence of operations?
 - $O(n)$
 - Why?
- Balanced Search Trees: stronger performance guarantees
- Main idea: rotation



BALANCED SEARCH TREES

- Single rotation: a constant number of parent-child relationships are modified
 - $O(1)$ for linked binary with a linked binary tree representation
- Rotations allow the shape of a tree to be modified while maintaining the search tree property
 - Rightward rotation: depth of each node in T_1 reduced by 1, depth of each node in T_3 increased by 1
- One or more rotation: trinode restructuring



BALANCED SEARCH TREES



Algorithm restructure(x):

Input: A position x of a binary search tree T that has both a parent y and a grandparent z

Output: Tree T after a trinode restructuring (which corresponds to a single or double rotation) involving positions x , y , and z

- 1: Let (a, b, c) be a left-to-right (inorder) listing of the positions x , y , and z , and let (T_1, T_2, T_3, T_4) be a left-to-right (inorder) listing of the four subtrees of x , y , and z not rooted at x , y , or z .
- 2: Replace the subtree rooted at z with a new subtree rooted at b .
- 3: Let a be the left child of b and let T_1 and T_2 be the left and right subtrees of a , respectively.
- 4: Let c be the right child of b and let T_3 and T_4 be the left and right subtrees of c , respectively.

BALANCED SEARCH TREES





Algorithm restructure(x):

Input: A position x of a binary search tree T that has both a parent y and a grandparent z

Output: Tree T after a trinode restructuring (which corresponds to a single or double rotation) involving positions x , y , and z

- 1: Let (a, b, c) be a left-to-right (inorder) listing of the positions x , y , and z , and let (T_1, T_2, T_3, T_4) be a left-to-right (inorder) listing of the four subtrees of x , y , and z not rooted at x , y , or z .
- 2: Replace the subtree rooted at z with a new subtree rooted at b .
- 3: Let a be the left child of b and let T_1 and T_2 be the left and right subtrees of a , respectively.
- 4: Let c be the right child of b and let T_3 and T_4 be the left and right subtrees of c , respectively.

BALANCED SEARCH TREES



Algorithm restructure(x):


Input: A position x of a binary search tree T that has both a parent y and a grandparent z

Output: Tree T after a trinode restructuring (which corresponds to a single or double rotation) involving positions x , y , and z

- 1: Let (a, b, c) be a left-to-right (inorder) listing of the positions x , y , and z , and let (T_1, T_2, T_3, T_4) be a left-to-right (inorder) listing of the four subtrees of x , y , and z not rooted at x , y , or z .
- 2: Replace the subtree rooted at z with a new subtree rooted at b .
- 3: Let a be the left child of b and let T_1 and T_2 be the left and right subtrees of a , respectively.
- 4: Let c be the right child of b and let T_3 and T_4 be the left and right subtrees of c , respectively.



AVL TREES

- AVL: Adelson-Velsky and Landis
- Adds a rule to the binary search tree to maintain a logarithmic height for the tree
- Height: number of edges on the longest path vs **number of nodes on this longest path**
 - Leaf position has height 1
- **Height balance property:** for every position p of T , the heights of the children of p differ by at most 1
- A subtree of an AVL tree is itself an AVL tree
- The height of an AVL tree storing n entries is $O(\log n)$
 - Proof in the text book





AVL TREES

- Insertion: insert item with key 54
- “search and repair”: going up from p to the root of T
 - z: first unbalanced position
 - y: child of z with higher height, y must be an ancestor of p
 - x: child of y with higher height (no tie, x must be an ancestor of p)
 - Call the trinode restructuring method, restructure(x)





AVL TREES

- Deletion: delete item with key 32
- Trinode restructuring:
 - z: first unbalanced position
 - y: child of z with larger height (y not ancestor of p)
 - x: child of y, such that if one the children of y is taller than the other, let x be the taller child of y. Else let x be the child of y on the same side as y
 - Perform restructure(x)
- Is this enough?



AVL TREES

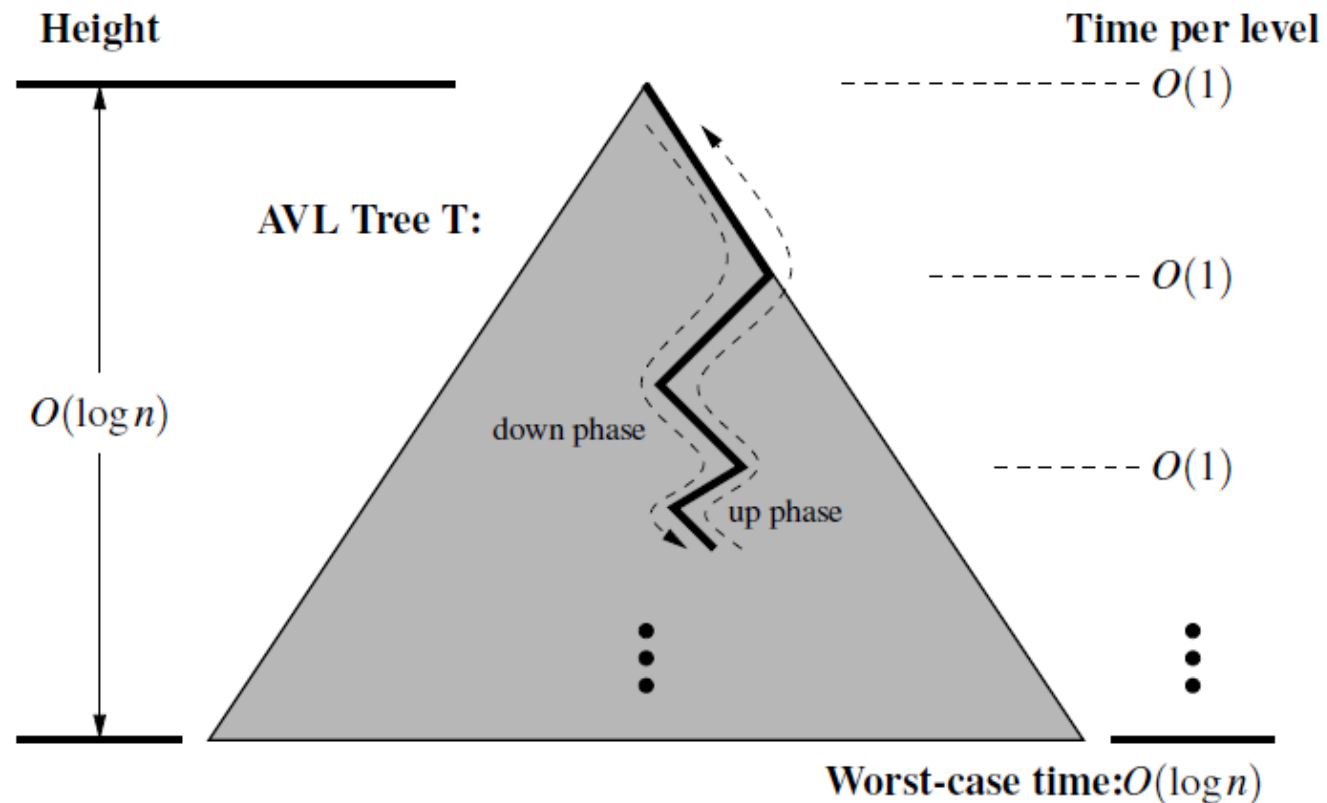
- Deletion: delete item with key 32
- Trinode restructuring:
 - May reduce the height of the subtree rooted at b by 1
 - Causes an ancestor of b to become unbalanced
 - Walk up T looking for unbalanced positions
 - $O(\log n)$ trinode restructuring are sufficient



PERFORMANCE OF AVL TREES

- AVL tree with n items: height guaranteed to be $O(\log n)$
- Standard BST operations: bounded by the height of tree
- AVL trees: $O(\log n)$ for most of the operations

Operation	Running Time
k in T	$O(\log n)$
$T[k] = v$	$O(\log n)$
$T.delete(p)$, $del\ T[k]$	$O(\log n)$
$T.find_position(k)$	$O(\log n)$
$T.first()$, $T.last()$, $T.find_min()$, $T.find_max()$	$O(\log n)$
$T.before(p)$, $T.after(p)$	$O(\log n)$
$T.find_lt(k)$, $T.find_le(k)$, $T.find_gt(k)$, $T.find_ge(k)$	$O(\log n)$
$T.find_range(start, stop)$	$O(s + \log n)$
$iter(T)$, $reversed(T)$	$O(n)$





SPLAY TREES(伸展树)

- Splay tree: different from the other balanced search trees
- No strict enforcement on a logarithmic upper bound on the height of the tree
- Efficiency realized by **splaying** operations
 - Performed at the bottommost position p reached for insertion, deletion, and search.
 - Splay operation causes more frequently accessed elements to remain nearer to the root
 - To reduce search times
 - Logarithmic amortised running time



SPLAY TREES(伸展树)

- Splay operations
- Given a node x of a binary search tree T , splay x – moving x to the root of T through a sequence of restructurings
- Zig-zig: node x and its parent y are both left children or both right children




SPLAY TREES(伸展树)

- Zig-zag: one of x and y is a left child and the other is a right child.
- Promote x by making x have y and z as its children, while maintaining the inorder relationships of the nodes in T




SPLAY TREES(伸展树)

- Zig: x does not have a grandparent
- Perform a single rotation to promote x over y making y a child of x




SPLAY TREES(伸展树)



(a)


Splaying the node storing 14 with a zig-zag



(b)


After the zig-zag

SPLAY TREES(伸展树)



(c)


Zig-zig




(d)

After the zig-zig

SPLAY TREES(伸展树)



Zig-zig




After zig-zig (f)

SPLAY TREES(伸展树)


- When to splay
 - When searching for key k , if k is found at position p , splay p ;
 - else splay the leaf position at which the search terminates unsuccessfully
 - When inserting key k , splay the newly created internal node where k gets inserted



(a)




(b)




(c)

SPLAY TREES(伸展树)


- When to splay
 - When searching for key k, if k is found at position p, splay p;
 - else splay the leaf position at which the search terminates unsuccessfully
 - When inserting key k, splay the newly created internal node where k gets inserted




(d)



(e)





(f)



(g)


SPLAY TREES(伸展树)

- When to splay
 - When deleting a key, splay the position p that is the parent of the removed node




SPLAY TREES(伸展树)

- When to splay
 - When deleting a key, splay the position p that is the parent of the removed node




(c)



(d)

SPLAY TREES(伸展树)

- When to splay
 - When deleting a key, splay the position p that is the parent of the removed node



(e)



THANKS

See you in the next session!