



GRAPHS

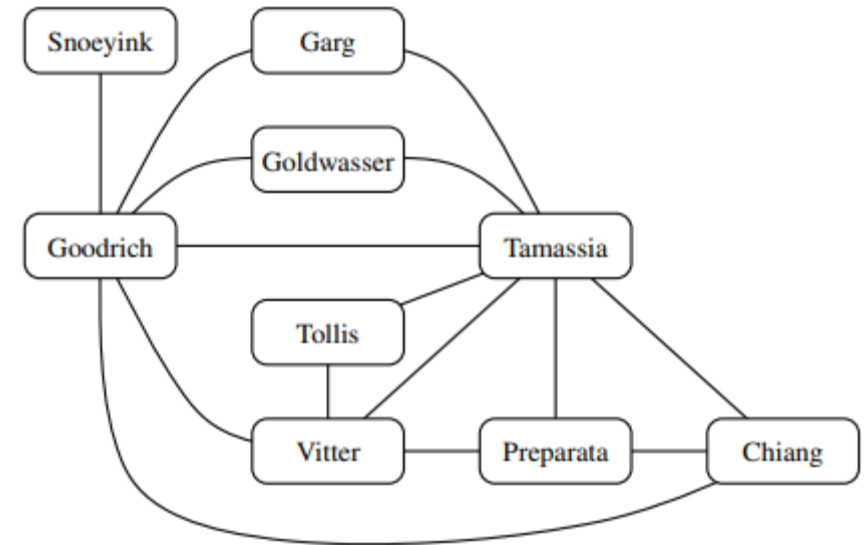
School of Artificial Intelligence

PREVIOUSLY ON GRAPH

- Graphs: relationships between pairs of objects: vertices together with a collection of connections between them (edges).
 - Mapping, transportation, computer networks electrical engineering
- Definition:
 - Set V of **vertices** and a collection E of pairs of vertices from V , called **edges**
- Edges
 - Directed: an edge (u, v) is directed from u to v if the pair (u, v) is ordered
 - Undirected: an edge (u, v) is not ordered, sometimes denoted as $\{u, v\}$
- Undirected graph: the edges are all undirected
- Directed graph (digraph): edges are all directed
- Mixed graph: directed edges + undirected edges
- Undirected graph \Rightarrow directed graph

GRAPHS

- Example
 - Collaboration graph (undirected)
 - Class relationships within an Object-Oriented program (directed)
 - City map (mixed graph)
 - Electric wiring/Schematics
 - Directed/undirected




GRAPHS

- **End vertices** (endpoints)
 - Two vertices joined by an edge
 - Directed edge: origin \rightarrow destination
- **Adjacent**(相邻的) vertices: u and v are adjacent if there is an edge whose end vertices are u and v
- **Incident**(入射): an edge is incident to a vertex if the vertex is one of the edge's endpoints
- **Outgoing**(输出) edges of a vertex: directed edges whose origin is the vertex
- **Incoming** (输入) edges of a vertex: directed edges whose destination is the vertex
- **Degree** (度) of a vertex: # of incident edges of v
- **In-degree** (入度) : # of incoming edges of v
- **Out-degree** (出度) : # of outgoing edges of v

GRAPHS

- End vertices (endpoints)
 - Two vertices joined by an edge
 - Directed edge: origin -> destination
- Adjacent vertices: u and v are adjacent if there is an edge whose end vertices are u and v
- Incident: an edge is incident to a vertex if the vertex is one of the edge's endpoints
- Outgoing edges of a vertex: directed edges whose origin is the vertex
- Incoming edges of a vertex: directed edges whose destination is the vertex
- Degree of a vertex: # of incident edges of v
- In-degree: # of incoming edges of v
- Out-degree: # of outgoing edges of v




GRAPHS

- Definition:
 - Set V of **vertices** and a collection E of pairs of vertices from V , called **edges**
- E : collection, not a set
 - Two edges can have the same end vertices
 - Such edges are called **parallel** or **multiple** edges
 - **Self-loop** edges: endpoints are the same
- **Simple** graphs: graphs do not have parallel edges or self loops
- **Path**: sequence of vertices and edges, start at a vertex and ends at a vertex
 - Simple: if each vertex in the path is distinct
- **Cycle**: path that starts and ends at the same vertex, and includes at least one edge
 - Simple: if each vertex in the cycle is distinct(except for the first and last one)
- **Acyclic**: directed graph that has no directed cycles

GRAPHS


- Directed path
 - A path such that all edges are directed and are traversed along their direction
 - (BOS, NW35, JFK, AA1387, DFW)
- Directed cycle
 - A cycle such that all edges are directed and are traversed along their direction
 - (LAX, UA1200, ORD, UA877, DFW, AA49, LAX)




GRAPHS

- Reachability
 - In a graph G , u **reaches** v , and v is reachable from u , if G has a path from u to v
- Connectivity
 - G is **connected**, if for any two vertices, there is a path between them
 - Directed graph: **strongly connected**
- Subgraph
 - Graph H whose vertices and edges are subsets of the vertices and edges of G
- Spanning subgraph
 - Subgraph of G that contains all the vertices of the graph
- Forest
 - Graph without cycles
- Tree
 - Connected forest
- Spanning tree
 - Spanning subgraph that is a tree


GRAPHS




(a)



(b)



(c)



(d)

GRAPHS

- If G is a graph with m edges and vertex set V , then
- If G is a directed graph with m edges and vertex set V , then
- Let G be a simple graph with n vertices and m edges.
 - If G is undirected, then $m \leq n(n-1)/2$
 - If G is directed, then $m \leq n(n-1)$
- Let G be a undirected graph with n vertices and m edges
 - If G is connected, then $m \geq n-1$
 - If G is a tree, then $m = n-1$
 - If G is a forest then $m \leq n-1$

$$\sum_{v \in V} \deg(v) = 2m.$$

$$\sum_{v \in V} \text{indeg}(v) = \sum_{v \in V} \text{outdeg}(v) = m.$$

GRAPHS

- Graph ADT
- Endpoints: return a tuple (u, v) such that vertex u is the origin and v the destination
- Opposite(v): if v is one endpoint of the edge, return the other endpoint
- Vertex_count(): returns # of vertices
- Vertices(): returns an iteration of all vertices of the graph
- Edge_count(): returns # of edges
- Edges(): returns an iteration of all edges
- Get_edge(u, v): returns the edge from u to v , if one exists, otherwise return None

GRAPHS

- Graph ADT
- `degree(v, out= True)`: returns the # of edges incident to v
- `incident_edges(v, out=True)`: returns an iteration of edges incident to v
- `insert_vertex(x = None)`: create and return a new Vertex storing element x
- `insert_edge(u, v, x = None)`: create and return a new Vertex from u to v
- `remove_vertex(v)`: remove v and all its incident edges from the graph
- `remove_edge(e)`: remove e from the graph

DATA STRUCTURES FOR GRAPHS


- Edge list
 - An unordered list of all edges, no efficient way to locate a particular edge (u, v) , or the set of all edges incident to a vertex v
- Adjacency list
 - For each vertex, a list containing edges that are incident to the vertex
 - Complete set of edges can be determined by taking the union of the smaller sets
- Adjacency map
 - Similar to adjacency list, but secondary container for edges are maps
 - $O(1)$ expected time to access a specific edge (u, v)
- Adjacency matrix
 - Worst case $O(1)$ access to a specific edge (u, v) by maintaining a $n \times n$ matrix for each graph with n vertices

DATA STRUCTURES FOR GRAPHS


Operation	Edge List	Adj. List	Adj. Map	Adj. Matrix
vertex_count()	$O(1)$	$O(1)$	$O(1)$	$O(1)$
edge_count()	$O(1)$	$O(1)$	$O(1)$	$O(1)$
vertices()	$O(n)$	$O(n)$	$O(n)$	$O(n)$
edges()	$O(m)$	$O(m)$	$O(m)$	$O(m)$
get_edge(u,v)	$O(m)$	$O(\min(d_u, d_v))$	$O(1)$ exp.	$O(1)$
degree(v)	$O(m)$	$O(1)$	$O(1)$	$O(n)$
incident_edges(v)	$O(m)$	$O(d_v)$	$O(d_v)$	$O(n)$
insert_vertex(x)	$O(1)$	$O(1)$	$O(1)$	$O(n^2)$
remove_vertex(v)	$O(m)$	$O(d_v)$	$O(d_v)$	$O(n^2)$
insert_edge(u,v,x)	$O(1)$	$O(1)$	$O(1)$ exp.	$O(1)$
remove_edge(e)	$O(1)$	$O(1)$	$O(1)$ exp.	$O(1)$

EDGE LIST

- Collections V and E are represented with doubly linked lists (PositionalList)
- Vertex objects
 - Reference to element x
 - Reference to position of the vertex instance in the list V
- Edge objects
 - Reference to element x
 - Reference to the vertex objects associated with the endpoint vertices of e
 - Reference to the position of the edge instance in list E



(a)



(b)


EDGE LIST

- Performance
- Space: $O(n + m)$
 - n vertices and m edges
- Running time
 - `vertices()`: $O(n)$
 - `edges()`: $O(m)$
 - `get_edge()`: $O(m)$
 - Most significant limitation
 - `remove_vertex(v)`: $O(m)$
 - Why?


Operation	Running Time
<code>vertex_count()</code> , <code>edge_count()</code>	$O(1)$
<code>vertices()</code>	$O(n)$
<code>edges()</code>	$O(m)$
<code>get_edge(u,v)</code> , <code>degree(v)</code> , <code>incident_edges(v)</code>	$O(m)$
<code>insert_vertex(x)</code> , <code>insert_edge(u,v,x)</code> , <code>remove_edge(e)</code>	$O(1)$
<code>remove_vertex(v)</code>	$O(m)$

ADJACENCY LIST

- Secondary containers for edges that are associated with each individual vertex
- For each v , maintain a collection $I(v)$ called **incidence collection** of v
- Primary structure: collection V of vertices
 - Positional list
- Each vertex instance
 - Direct reference to its $I(v)$ incidence collection



(a)



(b)


ADJACENCY LIST

- Performance
- Space: $O(n + m)$
 - n vertices and m edges
- Running time
 - `vertices()`: $O(n)$
 - `edges()`: $O(m)$
 - `get_edge()`: $O(\min(\deg(u), \deg(v)))$
 - Search through either $I(u)$ or $I(v)$
 - `remove_vertex(v)`: $O(\deg(v))$

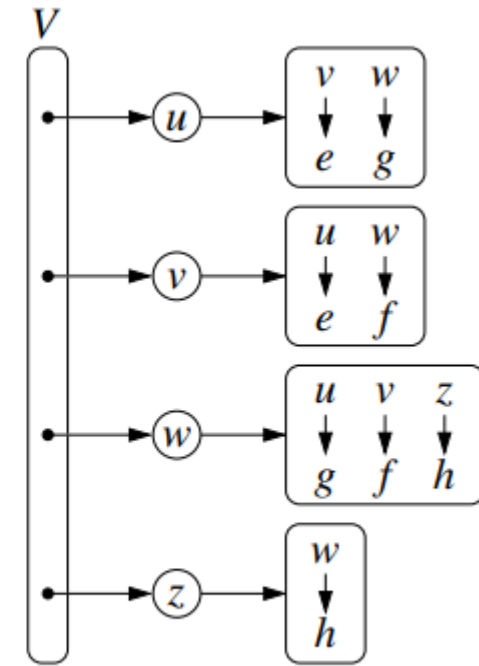
Operation	Running Time
<code>vertex_count()</code> , <code>edge_count()</code>	$O(1)$
<code>vertices()</code>	$O(n)$
<code>edges()</code>	$O(m)$
<code>get_edge(u,v)</code>	$O(\min(\deg(u), \deg(v)))$
<code>degree(v)</code>	$O(1)$
<code>incident_edges(v)</code>	$O(\deg(v))$
<code>insert_vertex(x)</code> , <code>insert_edge(u,v,x)</code>	$O(1)$
<code>remove_edge(e)</code>	$O(1)$
<code>remove_vertex(v)</code>	$O(\deg(v))$

ADJACENCY MAP

- Adjacency list
 - $I(v)$ uses $O(\deg(v))$ space
 - $O(\deg(v))$ time
- Performance improvement
 - Hash-based map for $I(v)$
 - $\text{get_edge}(u,v)$ can run in expected $O(1)$ time, worst case $O(\min(\deg(u), \deg(v)))$




(a)



(b)

ADJACENCY MATRIX

- Matrix A (n by n) to locate an edge between a given pair of vertices in worst-case $O(1)$ time
- Vertices as integers in $\{0, 1, \dots, n-1\}$
- $A[i, j]$ holds a reference to the edge (u, v) if one exists
- Edge (u, v) can be accessed in worst-case $O(1)$ time
- $O(n^2)$ space usage
- Matrix can be used to store only Boolean values, if edges do not store any additional data



(a)

		0	1	2	3
$u \rightarrow$	0		e	g	
$v \rightarrow$	1	e		f	
$w \rightarrow$	2	g	f		h
$z \rightarrow$	3			h	

(b)

THIS LECTURE: GRAPH TRAVERSALS

- Graph Traversals: a systematic procedure for exploring a graph by examining all of its vertices and edges
- Why traversal
 - Compute a path from u to v
 - Given a start vertex s of G , for every vertex v of G , compute the shortest path
 - Test whether G is connected
 - Compute a spanning tree of G if G is connected
 - Compute the connected components of G
 - Compute a cycle in G
 - Compute a directed path from u to v (for directed graphs)
 - Determine if G is acyclic (for directed graphs)
 - Determine if G is strongly connected (for directed graphs)

THIS LECTURE: GRAPH TRAVERSALS

- **Depth First Search**

- Maze solving
- Begin at a starting vertex s
 - s is the current vertex u
- For each edge of s
 - If it leads to a visited vertex, ignore
 - If it leads to an unvisited vertex v , make v the current vertex u , make u “visited”
 - repeat

Algorithm DFS(G, u): {We assume u has already been marked as visited}

Input: A graph G and a vertex u of G

Output: A collection of vertices reachable from u , with their discovery edges

for each outgoing edge $e = (u, v)$ of u **do**

if vertex v has not been visited **then**



 Mark vertex v as visited (via edge e).

 Recursively call DFS(G, v).


DFS

- **Depth First Search**
- Depth-first search tree: rooted at s
- When an edge $d(u, v)$ is used to discover a new vertex v , the edge is known as a discovery edge (**tree edge**)
- **Nontree edges**
 - **Back edges**: connect a vertex to an ancestor in the DFS tree
 - **Forward edges**: connect a vertex to a descendant in the DFS tree
 - **Cross edges**: connect a vertex to a vertex that is neither its ancestor nor its descendant


DFS




DFS




(a)




(b)




(e)



(f)



(c)



(d)

DFS

- **Depth First Search**

- G = undirected graph, when a DFS is performed at a starting vertex s , the traversal visits all vertices in the connected component of s , and the discovery edges form a spanning tree of the connected component of s
- G = directed graph, a DFS on G starting at vertex s visits all the vertices of G that are reachable from s . Also the DFS tree contains directed paths from s to every vertex reachable from s

- Running time

- $O(n_s + m_s)$
 - n_s : # of vertices reachable from a vertex s
 - m_s : # of incident edges to vertices reachable from s

DFS

- Implementation

```
1 def DFS(g, u, discovered):
2     """ Perform DFS of the undiscovered portion of Graph g starting at Vertex u.
3
4     discovered is a dictionary mapping each vertex to the edge that was used to
5     discover it during the DFS. (u should be "discovered" prior to the call.)
6     Newly discovered vertices will be added to the dictionary as a result.
7     """
8     for e in g.incident_edges(u):          # for every outgoing edge from u
9         v = e.opposite(u)
10        if v not in discovered:            # v is an unvisited vertex
11            discovered[v] = e              # e is the tree edge that discovered v
12            DFS(g, v, discovered)          # recursively explore from v
```

```
1 def construct_path(u, v, discovered):
2     path = [ ]
3     if v in discovered:
4         # we build list from v to u and then reverse it
5         path.append(v)
6         walk = v
7         while walk is not u:
8             e = discovered[walk]
9             parent = e.opposite(walk)
10            path.append(parent)
11            walk = parent
12        path.reverse( )
13    return path
```

DFS

- Applications of DFS
 - Testing for connectivity
 - Undirected graph, start from an arbitrary vertex and test if $\text{len}(\text{discovered}) == n$
 - Directed graph, test for strong connection $O(n(n+m))$ time, test for each (u, v) pairs
 - $O(n+m)$ is achieved by 2 DFSs, how?
 - Computing all connected components
 - Detecting cycles
 - Determine if there are **back edges**

```
1 def DFS_complete(g):
2     """ Perform DFS for entire
3
4     Result maps each vertex v
5     (Vertices that are roots of
6     """
7     forest = { }
8     for u in g.vertices():
9         if u not in forest:
10             forest[u] = None
11             DFS(g, u, forest)
12     return forest
```

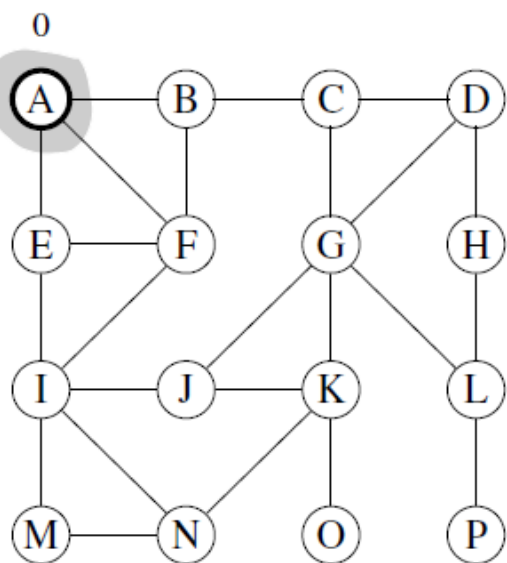
BFS

- **Breadth First Search**


- Web crawling
- Social network
- Garbage collection
- Rubic's cube
- Start from vertex s
- One 'level' at a time
- Level = all adjacent vertex from s ,
 - If vertex u is unvisited, mark u as 'visited'
 - If vertex u is visited, discard
- Proceed to the next 'level'

```
1 def BFS(g, s, discovered):
2     """ Perform BFS of the undiscovered vertices
3
4     discovered is a dictionary mapping vertices to discovery
5     discover it during the BFS (s should be discovered)
6     Newly discovered vertices will be added to the next level
7     """
8     level = [s]
9     while len(level) > 0:
10         next_level = []
11         for u in level:
12             for e in g.incident_edges(u):
13                 v = e.opposite(u)
14                 if v not in discovered:
15                     discovered[v] = e
16                     next_level.append(v)
17         level = next_level
```


BFS




(a)




(b)




(c)



(d)



(e)



(f)

BFS

- **Nontree edges**

- Undirected graph: all non tree edges are cross edges
- Directed graph: back edges or cross edges

- **Properties:**

- If G is a graph on which a BFS traversal starting at vertex s has been performed
 - The traversal visits all vertices of G that are reachable from s
 - For each vertex v at level i , the path of the BFS tree T between s and v has i edges and any other path of G from s to v has at least i edges
 - If (u, v) is an edge that is not in the BFS tree, then the level number of v can be at most 1 greater than the level number of u
- If G is a graph with n vertices and m edges, BFS traversal of G takes $O(n+m)$ time



TRANSITIVE CLOSURE

(传递闭包)

- Reachability problem in a directed graph
 - DFS or BFS: $O(n + m)$ time
- In certain applications, we want to answer many reachability queries more efficiently
- Transitive closure of a directed graph G
 - Is itself a directed graph G' , such that
 - The vertices of G' are the same as the vertices of G
 - G' has an edge (u, v) , when ever G has a directed path from u to v
- Computation of transitive closure $O(n(n+m))$
 - N DFSs
- Floyd-Warshall: $O(n^3)$ – not covered in this course

DIRECTED ACYCLIC GRAPHS

- Directed Acyclic Graph (DAG)
 - Directed graphs without directed cycles
- Applications
 - Prerequisites between courses of a degree program
 - Inheritance between classes of an object-oriented program
 - Scheduling constraints between the tasks of a project
- Topological Ordering
 - If G is a directed graph with n vertices
 - Topological ordering of G : an ordering v_1, v_2, \dots, v_n of vertices of G such that for every edge (v_i, v_j) of G , it is the case that $i < j$




```


1 def topological_sort(g):
2     """ Return a list of vertices of directed acy
3
4     If graph g has a cycle, the result will be inc
5     """
6     topo = [ ]           # a list of vertices
7     ready = [ ]          # list of vertices th
8     incount = { }        # keep track of in-
9     for u in g.vertices():
10         incount[u] = g.degree(u, False) # pa
11         if incount[u] == 0:             # if
12             ready.append(u)             # it
13     while len(ready) > 0:
14         u = ready.pop( )                 # u
15         topo.append(u)                   # ac
16         for e in g.incident_edges(u):    # co
17             v = e.opposite(u)
18             incount[v] -= 1               # v
19             if incount[v] == 0:
20                 ready.append(v)
21     return topo

```


DIRECTED AC



(a)



(b)






(c)

```

1 def topological_sort(g):
2     """ Return a list of vertices of directed a
3
4     If graph g has a cycle, the result will be
5     """
6     topo = []           # a list of vertices
7     ready = []          # list of vertices
8     incount = { }       # keep track of i
9     for u in g.vertices():
10         incount[u] = g.degree(u, False)
11         if incount[u] == 0:
12             ready.append(u)
13     while len(ready) > 0:
14         u = ready.pop( )
15         topo.append(u)
16         for e in g.incident_edges(u):
17             v = e.opposite(u)
18             incount[v] -= 1
19             if incount[v] == 0:
20                 ready.append(v)
21     return topo
    
```




DIRECTED AC



```

1 def topological_sort(g):
2     """ Return a list of vertices of directed
3
4     If graph g has a cycle, the result will be
5     """
6     topo = []           # a list of vertices
7     ready = []          # list of vertices with in-degree 0
8     incount = { }       # keep track of in-degrees
9     for u in g.vertices():
10        incount[u] = g.degree(u, False)
11        if incount[u] == 0:
12            ready.append(u)
13    while len(ready) > 0:
14        u = ready.pop()
15        topo.append(u)
16        for e in g.incident_edges(u):
17            v = e.opposite(u)
18            incount[v] -= 1
19            if incount[v] == 0:
20                ready.append(v)
21    return topo
    
```

DIRECTED AC



```

1 def topological_sort(g):
2     """ Return a list of vertices of directed
3
4     If graph g has a cycle, the result will be
5     """
6     topo = []           # a list of vertices
7     ready = []          # list of vertices with in-degree 0
8     incount = { }       # keep track of in-degrees
9     for u in g.vertices():
10         incount[u] = g.degree(u, False)
11         if incount[u] == 0:
12             ready.append(u)
13     while len(ready) > 0:
14         u = ready.pop()
15         topo.append(u)
16         for e in g.incident_edges(u):
17             v = e.opposite(u)
18             incount[v] -= 1
19             if incount[v] == 0:
20                 ready.append(v)
21     return topo
    
```




THANKS

See you in the next session!